

# *SIMD in C++*

Presented  
January, 2017



BrightSource

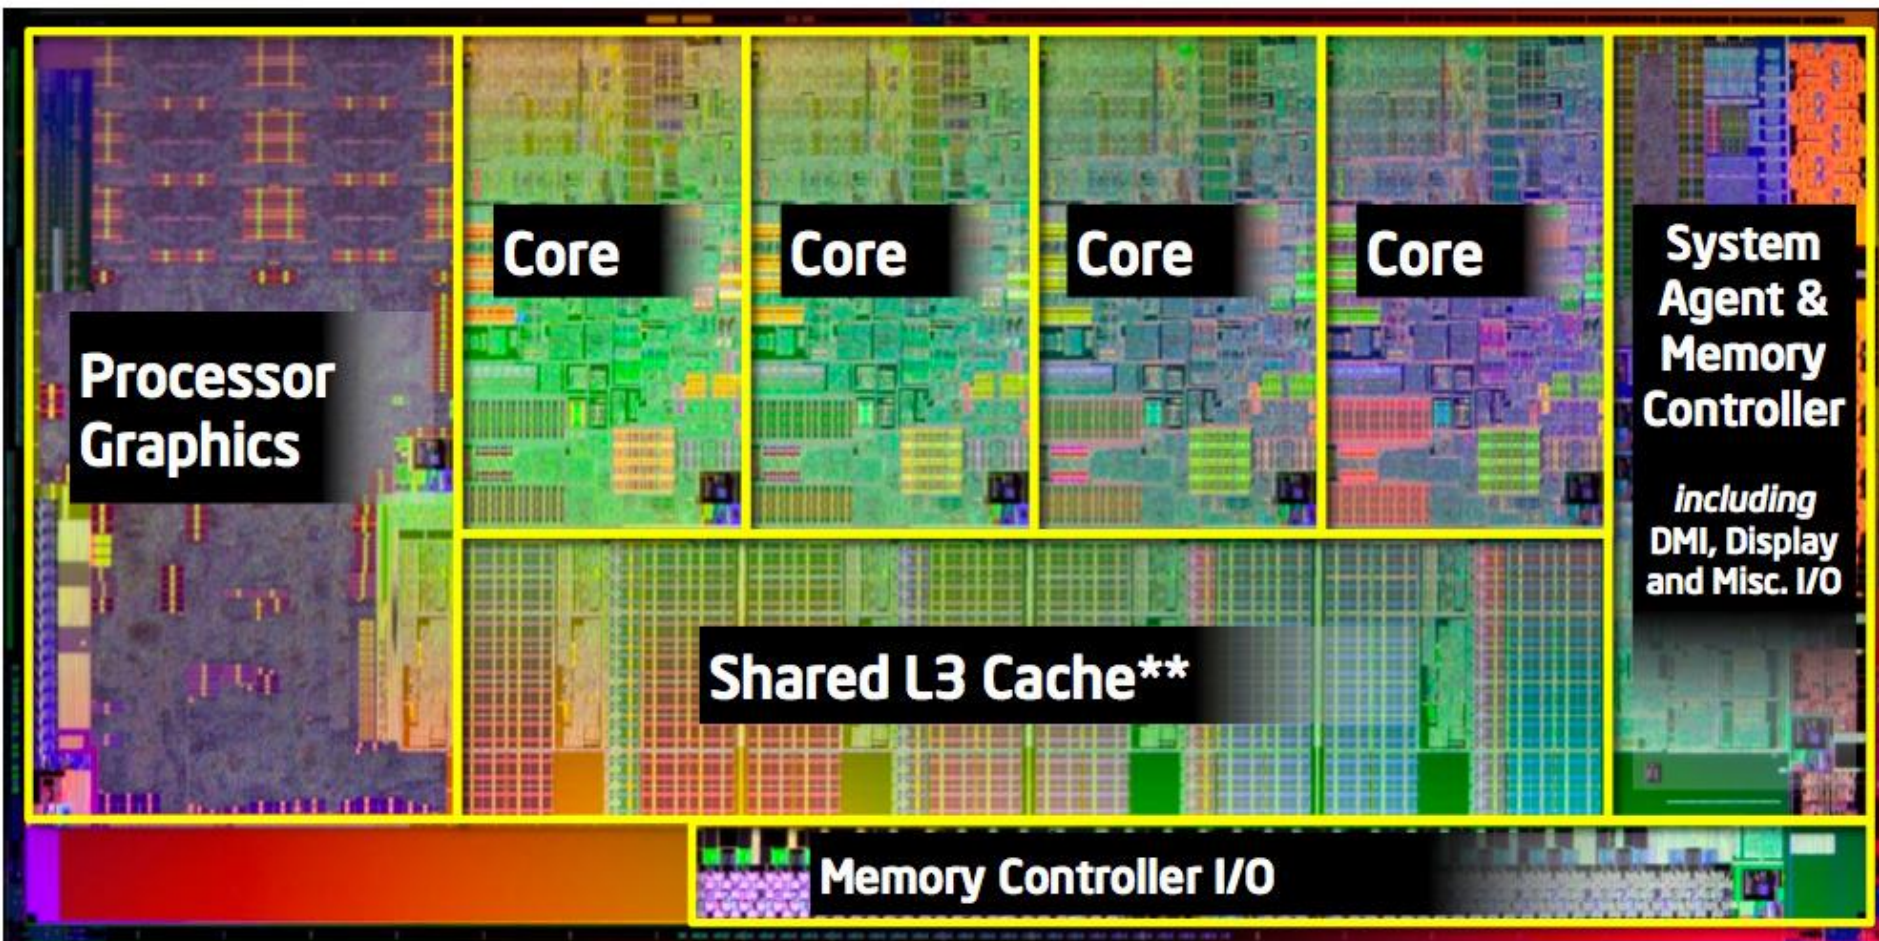
Doc ID:

© 2006-2017 BrightSource Energy, Inc. and/or any of its subsidiaries (BrightSource)  
The contents of this document and any equipment and/or software described may be  
protected by trade secrets and patents and patents pending in the People's Republic  
of China, the United States and/or other countries. This document contains trade  
secrets and other information that is confidential and proprietary to BrightSource. Any  
reproduction, disclosure or other use thereof is prohibited except as otherwise agreed  
to in advance and in writing by BrightSource.





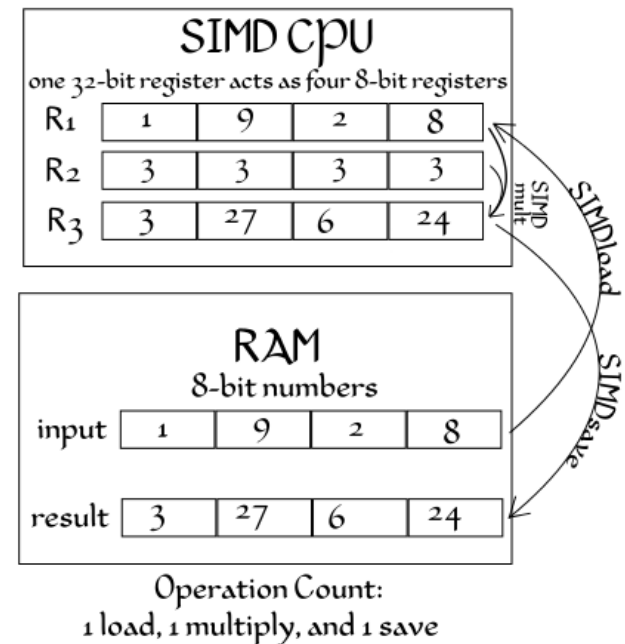
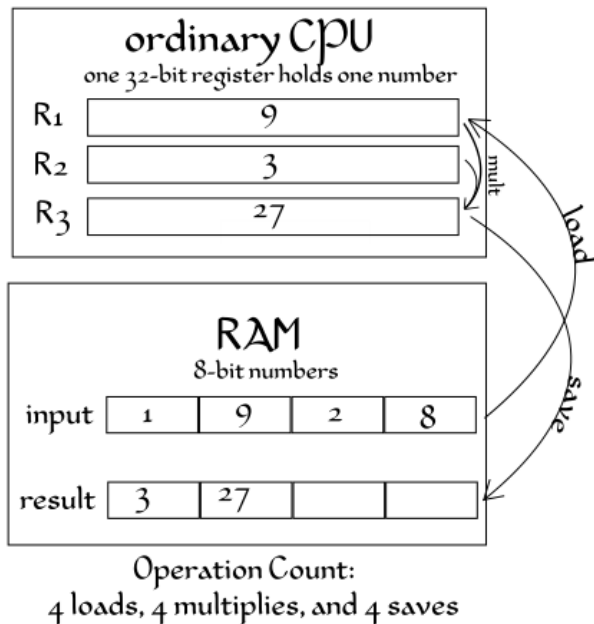


# Why C++?







# What is SIMD





# History

- Intel
  - SSE – Streaming SIMD Extensions
  - 128 bit registers, XMM
    - 1999 – Version1
    - 2001 – Version2
    - 2003 – Version3
    - 2006 – Version4
  - AVX - Advanced Vector Extensions
  - 256 bit registers, YMM
    - 2011 – Version1
    - 2013 – Version2
  - AVX-512
    - 2015 – AVX 512
  - ...

AVX-512 register scheme as extension from the AVX (YMM0-YMM15) and SSE (XMM0-XMM15) registers

511	256	255	128	127	0
ZMM0	YMM0	XMM0			
ZMM1	YMM1	XMM1			
ZMM2	YMM2	XMM2			
ZMM3	YMM3	XMM3			
ZMM4	YMM4	XMM4			
ZMM5	YMM5	XMM5			
ZMM6	YMM6	XMM6			
ZMM7	YMM7	XMM7			
ZMM8	YMM8	XMM8			
ZMM9	YMM9	XMM9			
ZMM10	YMM10	XMM10			
ZMM11	YMM11	XMM11			
ZMM12	YMM12	XMM12			
ZMM13	YMM13	XMM13			
ZMM14	YMM14	XMM14			
ZMM15	YMM15	XMM15			
ZMM16	YMM16	XMM16			
ZMM17	YMM17	XMM17			
ZMM18	YMM18	XMM18			
ZMM19	YMM19	XMM19			
ZMM20	YMM20	XMM20			
ZMM21	YMM21	XMM21			
ZMM22	YMM22	XMM22			
ZMM23	YMM23	XMM23			
ZMM24	YMM24	XMM24			
ZMM25	YMM25	XMM25			
ZMM26	YMM26	XMM26			
ZMM27	YMM27	XMM27			
ZMM28	YMM28	XMM28			
ZMM29	YMM29	XMM29			
ZMM30	YMM30	XMM30			
ZMM31	YMM31	XMM31			







# Issues

## ■ Data alignment

- For optimal performance SIMD data needs to be aligned in memory
- In order to load to a register of size 128, the data needs to be aligned to 128
- It's possible to load/store unaligned data, you might pay performance penalty, depending on hardware
- <http://en.cppreference.com/w/cpp/language/alignof>
- <http://en.cppreference.com/w/cpp/language/alignas>








# Issues - continued

- Data arranging, reorg
  - In order for SIMD to be efficient you might want to arrange the data “by properties”
  - Not so good (Array of Structures):
    - `std::vector<Point3> points`
  - Better (Structure of Arrays)
    - `struct Points { std::vector<float> Xs, Ys, Zs; }`
  - Though there are operations you can do within the register...



Point Structure with Data in AoS Arrangement

```
struct Point{  
    float r;  
    float g;  
    float b;  
}
```



Points Structure with Data in SoA Arrangement

```
struct Points{  
    float* x;  
    float* y;  
    float* z;  
}
```





# Don't do it yourself

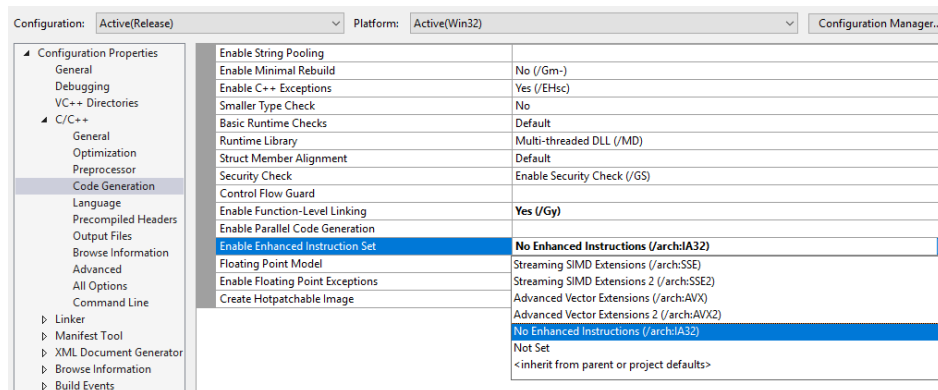
- We have libraries!
  - Some are Low level:
    - <https://github.com/NumScale/boost.simd>
    - <http://ermig1979.github.io/Simd/help/index.html>
    - <https://software.intel.com/en-us/intel-ipp>
  - Some are High level
    - <http://eigen.tuxfamily.org/>
    - <https://opencv.org/>
  - Compilers usually do some of it for you, familiarize yourself with it:

```
#include <boost/simd/pack.hpp>
#include <iostream>

namespace bs = boost::simd;

int main()
{
    bs::pack<float,4> p{1.f,2.f,3.f,4.f};
    std::cout << p + 10*p << "\n";

    return 0;
}
```





# Is your software compute or data bound?

- <http://www.overbyte.com.au/misc/Lesson3/CacheFun.html>
- <https://gist.github.com/jboner/2841832>

## Latency Comparison Numbers

-----

L1 cache reference	0.5	ns			
Branch mispredict	5	ns			
L2 cache reference	7	ns			14x L1 cache
Mutex lock/unlock	25	ns			
Main memory reference	100	ns			20x L2 cache, 200x L1 cache
Compress 1K bytes with Zippy	3,000	ns	3	us	
Send 1K bytes over 1 Gbps network	10,000	ns	10	us	
Read 4K randomly from SSD*	150,000	ns	150	us	~1GB/sec SSD
Read 1 MB sequentially from memory	250,000	ns	250	us	
Round trip within same datacenter	500,000	ns	500	us	
Read 1 MB sequentially from SSD*	1,000,000	ns	1,000	us	1 ms ~1GB/sec SSD, 4X memory
Disk seek	10,000,000	ns	10,000	us	10 ms 20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	20,000	us	20 ms 80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000	ns	150,000	us	150 ms



# GP-GPU Parallelization Alternative

- Another option to gain performance on dedicated hardware is using the GPU processor abilities for General Purpose

```
#include <vector>
#include <algorithm>
#include <boost/compute.hpp>

namespace compute = boost::compute;

int main()
{
    // get the default compute device
    compute::device gpu = compute::system::default_device();

    // create a compute context and command queue
    compute::context ctx(gpu);
    compute::command_queue queue(ctx, gpu);

    // generate random numbers on the host
    std::vector<float> host_vector(1000000);
    std::generate(host_vector.begin(), host_vector.end(), rand);

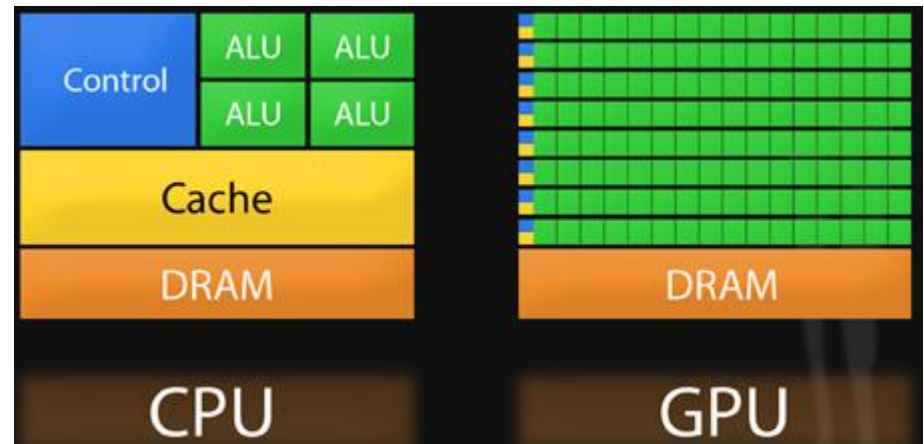
    // create vector on the device
    compute::vector<float> device_vector(1000000, ctx);

    // copy data to the device
    compute::copy(
        host_vector.begin(), host_vector.end(), device_vector.begin(), queue
    );

    // sort data on the device
    compute::sort(
        device_vector.begin(), device_vector.end(), queue
    );

    // copy data back to the host
    compute::copy(
        device_vector.begin(), device_vector.end(), host_vector.begin(), queue
    );

    return 0;
}
```



- Very low level, separate area of expertise
- High latency for data transfer to/from GPU
- Lots of primitive CPUs
- <https://github.com/boostorg/compute>



**BrightSource**  
*Limitless*

[brightsourceenergy.com](http://brightsourceenergy.com)

