

HFT

The Old days



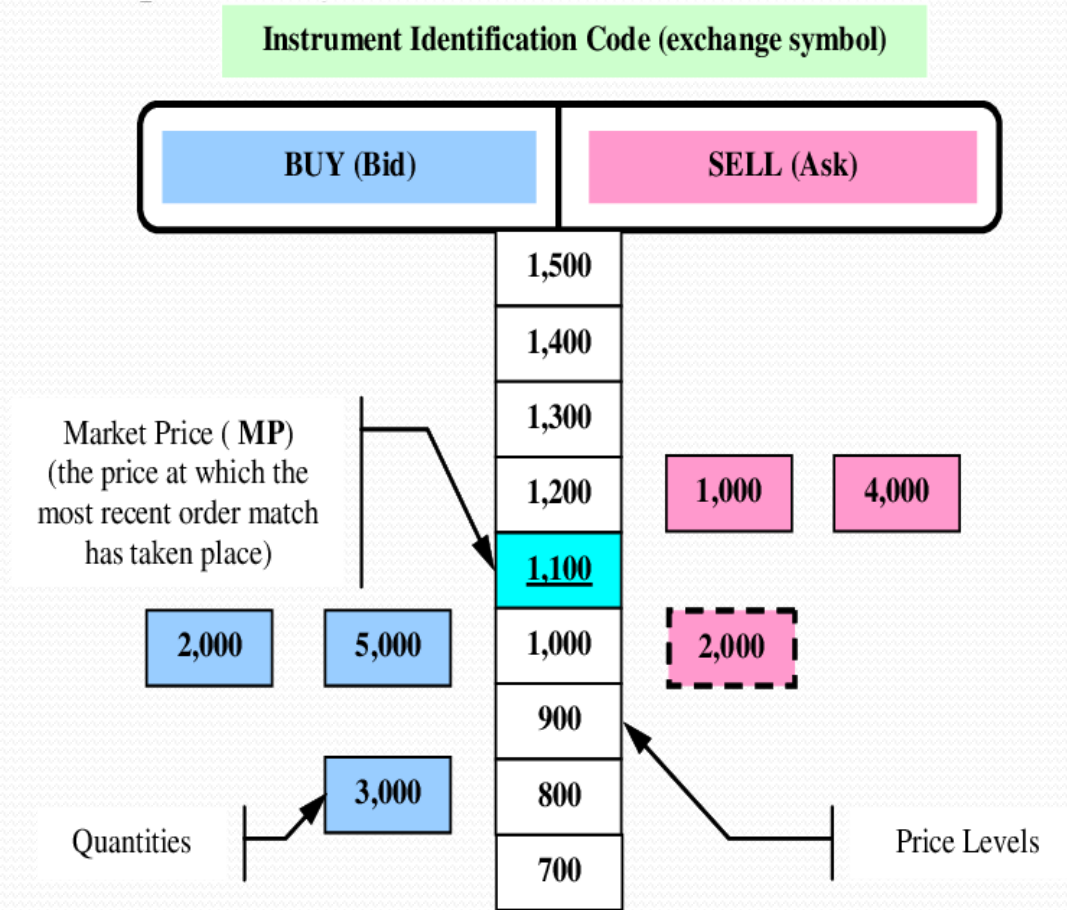
- Trading floor – CBOT 2006

Today



- Trading BlackBerry stock - 2013

Basic Order flow



How to be fast

- Physical location
- Know your network
- Know your OS
- Measure, Measure, Measure
- Write fast code

Write fast code

```
int main()
{
    set_priority_and_affinity();
    // Generate data
    const unsigned arraySize = 32768;
    int data[arraySize];

    for (unsigned c = 0; c < arraySize; ++c)
        data[c] = std::rand() % 256;

    // Test
    clock_t start = clock();

    long long sum = 0;

    for (unsigned i = 0; i < 100000; ++i)
    {
        // Primary loop
        for (unsigned c = 0; c < arraySize; ++c)
        {
            if (data[c] >= 128)
                sum += data[c];
        }
    }

    double elapsedTime = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;

    std::cout << elapsedTime << std::endl;
    std::cout << "sum = " << sum << std::endl;
    return 0;
}
```


Write fast code

```
int main()
{
    set_priority_and_affinity();
    // Generate data
    const unsigned arraySize = 32768;
    int data[arraySize];

    for (unsigned c = 0; c < arraySize; ++c)
        data[c] = std::rand() % 256;

    // Test
    clock_t start = clock();

    std::sort(data, data + arraySize);

    long long sum = 0;
    for (unsigned i = 0; i < 100000; ++i)
    {
        // Primary loop
        for (unsigned c = 0; c < arraySize; ++c)
        {
            if (data[c] >= 128)
                sum += data[c];
        }
    }

    double elapsedTime = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;
    std::cout << elapsedTime << std::endl;
    std::cout << "sum = " << sum << std::endl;
    return 0;
}
```

```
# perf stat ./branch_test_unsorted_noop
```

```
22.4591
```

```
sum = 314931600000
```

```
Performance counter stats for './branch_test_unsorted_noop':
```

22460.791733	task-clock (msec)	#	1.000 CPUs utilized
9	context-switches	#	0.000 K/sec
1	cpu-migrations	#	0.000 K/sec
359	page-faults	#	0.016 K/sec
62887815278	cycles	#	2.800 GHz
16427111065	stalled-cycles-frontend	#	26.12% frontend cycles idle
13767229572	stalled-cycles-backend	#	21.89% backend cycles idle
32831495807	instructions	#	0.52 insns per cycle
		#	0.50 stalled cycles per insn
9836478981	branches	#	437.940 M/sec
1543738105	branch-misses	#	15.69% of all branches

```
22.460951456 seconds time elapsed
```

```
# perf stat ./branch_test_sorted_noop
```

```
8.52512
```


```
sum = 314931600000
```

```
Performance counter stats for './branch_test_sorted_noop':
```

8532.642974	task-clock (msec)	#	1.000 CPUs utilized
3	context-switches	#	0.000 K/sec
1	cpu-migrations	#	0.000 K/sec
263	page-faults	#	0.031 K/sec
23890488385	cycles	#	2.800 GHz
14020983829	stalled-cycles-frontend	#	58.69% frontend cycles idle
493056149	stalled-cycles-backend	#	2.06% backend cycles idle
32835577975	instructions	#	1.37 insns per cycle
		#	0.43 stalled cycles per insn
9837162715	branches	#	1152.886 M/sec
351634	branch-misses	#	0.00% of all branches

```
8.532841461 seconds time elapsed
```


CPU pipeline



```
#define ROW 256
#define COL 256
int main(int argc, char** argv)
{
    int matrix[ROW][COL];
    int i,j;
    double elapsedTime;
    long long sum=0;
    clock_t start;

    set_priority_and_affinity();
    for (i=0; i<ROW;i++)
        for(j=0;j<COL;j++)
            matrix[i][j] = std::rand() % 256;

    start = clock();
    for (unsigned k = 0; k < 100000; ++k)
    {
        for (i=0; i<ROW;i++)
            for(j=0;j<COL;j++)
                sum += matrix[i][j];
    }
    elapsedTime = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;

    std::cout << elapsedTime << std::endl;
    std::cout << "sum = " << sum << std::endl;
}
```

```
#define ROW 256
#define COL 256
int main(int argc, char** argv)
{
    int matrix[ROW][COL];
    int i,j;
    double elapsedTime;
    long long sum=0;
    clock_t start;

    set_priority_and_affinity();
    for (i=0; i<ROW;i++)
        for(j=0;j<COL;j++)
            matrix[i][j] = std::rand() % 256;

    start = clock();
    for (unsigned k = 0; k < 100000; ++k)
    {
        for (i=0; i<COL;i++)
            for(j=0;j<ROW;j++)
                sum += matrix[j][i];
    }
    elapsedTime = static_cast<double>(clock() - start) / CLOCKS_PER_SEC;

    std::cout << elapsedTime << std::endl;
    std::cout << "sum = " << sum << std::endl;
}
```

```
# perf stat -d ./cache_test_cache 1
```

```
17.7906
```

```
sum = 837443300000
```

```
Performance counter stats for './cache_test_cache 1':
```

17792.660394	task-clock (msec)	#	1.000 CPUs utilized
6	context-switches	#	0.000 K/sec
1	cpu-migrations	#	0.000 K/sec
450	page-faults	#	0.025 K/sec
49817551332	cycles	#	2.800 GHz
22649652231	stalled-cycles-frontend	#	45.47% frontend cycles idle
774349774	stalled-cycles-backend	#	1.55% backend cycles idle
85405384703	instructions	#	1.71 insns per cycle
		#	0.27 stalled cycles per insn
13189583656	branches	#	741.294 M/sec
25737998	branch-misses	#	0.20% of all branches
39407406531	L1-dcache-loads	#	2214.812 M/sec
410719902	L1-dcache-load-misses	#	1.04% of all L1-dcache hits
11091380	LLC-loads	#	0.623 M/sec
12424	LLC-load-misses	#	0.11% of all LL-cache hits

```
17.792681123 seconds time elapsed
```

```
# perf stat -d ./cache_test_cache 0
```

```
19.4345
```

```
sum = 837443300000
```

```
Performance counter stats for './cache_test_cache 0':
```

19436.644440	task-clock (msec)	#	1.000 CPUs utilized
8	context-switches	#	0.000 K/sec
1	cpu-migrations	#	0.000 K/sec
518	page-faults	#	0.027 K/sec
54420463774	cycles	#	2.800 GHz
27599861890	stalled-cycles-frontend	#	50.72% frontend cycles idle
1235962601	stalled-cycles-backend	#	2.27% backend cycles idle
85408175572	instructions	#	1.57 insns per cycle
		#	0.32 stalled cycles per insn
13190066445	branches	#	678.618 M/sec
25733621	branch-misses	#	0.20% of all branches
39408260709	L1-dcache-loads	#	2027.524 M/sec
6562156690	L1-dcache-load-misses	#	16.65% of all L1-dcache hits
3335789969	LLC-loads	#	171.624 M/sec
20409	LLC-load-misses	#	0.00% of all LL-cache hits

```
19.436667480 seconds time elapsed
```

Code challenge



Write fast code – real example

- MD prices are published by the exchange as scaled integer, we need to generate code which changes the exponent as fast as possible

Field Name	Type	Values
<u>msg_type</u>	uint8	0x8000 'P' - snapshot 0x0000 'P' - incremental
<u>side_and_index</u>	uint8	MSB – side (1 – bid, 0 – ask) Rest of the bits – starting level in book
count	uint8	Number of units
<u>Presence_map</u>	uint8	See above
Units:		
price	int64	Price in this level
<u>distance_from_best</u>	uint16	Distance in ticks from best price
<u>num_of_orders</u>	uint16	Number of orders that comprise this price level
size	uint32	Quantity

Write fast code – real example

True Value	Equalsto ...	Equals to...
1,500	$1500 * 10^0$	$15 * 10^2$
1,500,000,000	$150000 * 10^4$	$1500 * 10^6$
0.15	$15 * 10^{-2}$	$1500 * 10^{-4}$

Base Version (66ns)

```
boost::int64_t changeExponent(boost::int64_t mantissa, boost::int16_t from_exp,
                              boost::int16_t to_exp)
{
    boost::int16_t diff = to_exp - from_exp;
    if (!diff)
        return mantissa;
    if (diff <= 0)
    {
        boost::int64_t div = (boost::int64_t)pow(10.0, (double)abs(diff));
        return mantissa * div;
    }
    else
    {
        boost::uint32_t div = (boost::uint32_t)pow(10.0, (double)abs(diff));
        return mantissa / div;
    }
}
```

Benchmark utility

```
bench_result_t benchmark_candidate(
    validation_input_t inputs,
    int warmup_iterations,
    int num_iterations,
    CandidateContainer& candidate
)
{
    flush_cache();
    for (int i = 0; i < warmup_iterations; i++)
    {
        const auto& cur_input = inputs[i];
        const auto& expected_result = std::get<0>(cur_input);
        const auto& func_params = std::get<1>(cur_input);
        auto func_ptr = candidate.get_func_ptr();
        auto mantissa = func_ptr(std::get<0>(func_params), std::get<1>(func_params), std::get<2>(func_params));
        dummy_result += mantissa;
    }
    boost::uint64_t run_result = 0;
    double run_time_ns;
    boost::uint64_t failed = 0;
    boost::uint32_t begin_high, begin_low, end_high, end_low;
    auto start = std::chrono::high_resolution_clock::now();
    for (int i = 0; i < num_iterations; i++)
    {
        const auto& cur_input = inputs[i];
        const auto& expected_result = std::get<0>(cur_input);
        const auto& func_params = std::get<1>(cur_input);
        auto func_ptr = candidate.get_func_ptr();

        auto orig_mantissa = std::get<0>(func_params);
        auto from_exp = std::get<1>(func_params);
        auto to_exp = std::get<2>(func_params);

        auto mantissa = func_ptr(orig_mantissa, from_exp, to_exp);
        run_result ^= mantissa;

        if (mantissa != expected_result)
            failed++;
    }
    auto diff = std::chrono::high_resolution_clock::now() - start;
    run_time_ns = (double)std::chrono::duration_cast<std::chrono::nanoseconds>(diff).count();
    run_time_ns /= num_iterations;
    return std::make_tuple(run_time_ns, run_result, failed);
}
```

Solution 1 :

```
static boost::int64_t MUL_POWERS_OF_10_BASE[] =
{
    0, // 0
    0,
    0,
    0,
    0,
    0, // 5
    0,
    0,
    0,
    0,

    10,
    100,
    1000,
    10000,
    100000,
    1000000,
    10000000,
    100000000,
    1000000000,
    10000000000
};
```

[illegible]

Solution 1:

```
static boost::int64_t* MUL_POWERS_OF_10 = &MUL_POWERS_OF_10_BASE[9];
static boost::int64_t* DIV_POWERS_OF_10 = &DIV_POWERS_OF_10_BASE[9];

boost::int64_t solution1(boost::int64_t mantissa, boost::int16_t
    from_exp, boost::int16_t to_exp)
{
    boost::int16_t diff = from_exp - to_exp;
    return mantissa / DIV_POWERS_OF_10[diff] + mantissa *
        MUL_POWERS_OF_10[diff];
}
```

Solution 1_(28ns):

```
sub    %edx,%esi
mov    %rdi,%rax
movswq %si,%rsi
cqto
imul   0x623428(,%rsi,8),%rdi
idivq  0x623388(,%rsi,8)
add    %rdi,%rax
retq
```


Solution 2:

```
boost::int64_t solution2(boost::int64_t mantissa, boost::int16_t from_exp, boost::int16_t to_exp)
{
    boost::int16_t diff = to_exp - from_exp;
    if (diff <= 0)
    {
        while (diff < -3) {
            mantissa *= 10000;
            diff += 4;
        }
        if (diff < -1) {
            mantissa *= 100;
            diff += 2;
        }
        if (diff < 0)
        {
            mantissa *= 10;
        }
    }
    else
    {
        while (diff > 3) {
            mantissa *= 0.0001;
            diff -= 4;
        }
        if (diff > 1) {
            mantissa *= 0.01;
            diff -= 2;
        }
        if (diff > 0)
        {
            mantissa *= 0.1;
        }
    }
    return mantissa;
}
```

Solution 2 (23ns):

```
sub %esi,%edx
mov %rdi,%rax
test %dx,%dx
mov %edx,%ecx
jle 0x40abd2 <uris(long, short, short)+146>
cmp $0x3,%dx
movsd 0x12554(%rip),%xmm1 # 0x41d0b0
jle 0x40ab86 <uris(long, short, short)+70>
xchg %ax,%ax
pxor %xmm0,%xmm0
sub $0x4,%ecx
cmp $0x3,%cx
cvtsi2sd %rax,%xmm0
mulsd %xmm1,%xmm0
cvttsd2si %xmm0,%rax
ja 0x40ab60 <uris(long, short, short)+32>
lea -0x4(%rdx),%ecx
and $0xffffffff,%ecx
sub %ecx,%edx
lea -0x4(%rdx),%ecx
cmp $0x1,%cx
jle 0x40aba5 <uris(long, short, short)+101>
pxor %xmm0,%xmm0
sub $0x2,%ecx
cvtsi2sd %rax,%xmm0
mulsd 0x12518(%rip),%xmm0 # 0x41d0b8
cvttsd2si %xmm0,%rax
cmp $0x1,%cx
jne 0x40abc1 <uris(long, short, short)+129>
pxor %xmm0,%xmm0
```

```
cvtsi2sd %rax,%xmm0
mulsd 0x12504(%rip),%xmm0 # 0x41d0c0
cvttsd2si %xmm0,%rax
repz retq
nopl 0x0(%rax,%rax,1)
imul $0x2710,%rax,%rax
add $0x4,%ecx
cmp $0xffff,%cx
jl 0x40abc8 <uris(long, short, short)+136>
cmp $0xffff,%cx
jl 0x40abec <uris(long, short, short)+172>
cmp $0xffff,%cx
jne 0x40abc1 <uris(long, short, short)+129>
lea (%rax,%rax,4),%rax
add %rax,%rax
retq
lea (%rax,%rax,4),%rax
add $0x2,%ecx
lea (%rax,%rax,4),%rax
shl $0x2,%rax
jmp 0x40abde <uris(long, short, short)+158>
```

Solution 3:

```
typedef boost::int64_t (*CalcFunc)(boost::int64_t mantissa);  
template<boost::int32_t X> boost::int64_t fastMul(boost::int64_t mantissa) { return mantissa * X; }  
template<boost::int32_t X> boost::int64_t fastDiv(boost::int64_t mantissa) { return mantissa / X; }  
boost::int64_t fastNone(boost::int64_t mantissa) { return mantissa; }
```

```
CalcFunc m_fastPow[31];  
void init()  
{  
    m_fastPow[0] = &fastMul<1000000000000000>;  
    m_fastPow[1] = &fastMul<1000000000000000>;  
    .  
    .  
    m_fastPow[15] = &fastNone;  
    m_fastPow[16] = &fastDiv<10>;  
    m_fastPow[17] = &fastDiv<100>;  
    .  
    .  
    m_fastPow[30] = &fastDiv<1000000000000000>;  
}
```

```
boost::int64_t solution3(boost::int64_t mantissa, boost::int16_t from_exp, boost::int16_t to_exp)  
{  
    return m_fastPow[to_exp - from_exp + 15](mantissa);  
}
```

Solution 3 (20ns):

```
movswl  %dx,%eax  
movswl  %si,%esi  
sub     %esi,%eax  
add     $0xf,%eax  
cltq  
mov     0x623e80(,%rax,8),%rax  
impg    *%rax
```

fastDiv<1000000000>(long):

```
mov      %rdi,%rax  
movabs   $0x112e0be826d694b3,%rdx  
sar      $0x3f,%rdi  
imul     %rdx  
sar      $0x1a,%rdx  
mov      %rdx,%rax  
sub      %rdi,%rax  
retq
```

Solution 4:

```
double PowArray4[] = {
    0.00000000000000000001,
    0.00000000000000000001,
    .
    .
    0.1,
    1L,
    10L,
    .
    .
    10000000000000000000L };

double* midPtr = &(PowArray4[19]);
signed long solution4(signed long mantissa, signed short from,
    signed short to)
{
    return mantissa * (*(midPtr + (from - to)));
}
```


Solution 4 (16.8ns):

```
pxor          %xmm0,%xmm0
mov          0x2173e5(%rip),%rax  # 0x622590 <midPtr>
movswl       %dx,%edx
movswl       %si,%esi
sub          %edx,%esi
cvtssi2sd    %rdi,%xmm0
movslq       %esi,%rsi
mulsd        (%rax,%rsi,8),%xmm0
cvttsd2si    %xmm0,%rax
retq
```


Solution 5:

```
static double exps_[] = {  
    0.0000000000000001,  
    0.0000000000000001,  
    .  
    .  
    0.1,  
    1,  
    10,  
    .  
    .  
    10000000000000000  
};  
static double* exps = exps_ + 15;  
  
inline int64_t solution5(int64_t m, int16_t from, int16_t to)  
{  
    return m * exps[from - to];  
}
```

Solution 5 (16.4ns):

<u>pxor</u>	%xmm0,%xmm0
<u>movswl</u>	% <u>dx</u> ,% <u>edx</u>
<u>movswl</u>	% <u>si</u> ,% <u>esi</u>
sub	% <u>edx</u> ,% <u>esi</u>
cvttsd2sd	%rdi,%xmm0
<u>movslq</u>	% <u>esi</u> ,% <u>rsi</u>
<u>mulsd</u>	0x623758(,%rsi,8),%xmm0
cvttsd2si	%xmm0,%rax
<u>retq</u>	

Solution 6:

```
static double exps_[] = {
    0.0000000000000001,
    0.0000000000000001,
    .
    .
    0.1,
    1,
    10,
    .
    .
    10000000000000000
};
static double* exps = exps_ + 15;

inline int64_t solution5(int64_t m, int16_t from, int16_t to)
{
    int16_t diff = from - to;
    return m * exps[diff];
}
```

Solution 6 (16 ns):

```
pxor          %xmm0,%xmm0
sub            %edx,%esi
movswq        %si,%rsi
cvtsi2sd      %rdi,%xmm0
mulsd         0x41a618(,%rsi,8),%xmm0
cvttsd2si     %xmm0,%rax
retq
```

Solution 7:

$$\frac{X}{1000} = \left(X \cdot \frac{2^{32}}{1000} \right) \div 2^{32} = (X \cdot 4294967.296) \div 2^{32} \\ \approx (X \cdot 4294967) \div 2^{32} = (X \cdot 4294967) \gg 32$$

General case:

$$\frac{X}{C} = \left(X \cdot \left\lfloor \frac{2^k}{C} \right\rfloor \right) \gg K$$

Error: $\frac{X}{2^k}$

Solution 7:

```
int64_t mult_val[41] = {0, .. , 1000 , 100, 10, 1, 107374182, 10737418,  
    1073741, .. , 107, 11,1,0,...,0 };  
uint64_t shift_val_array[41] = { 0,0,...,0,0,30,30,...,30 };  
int64_t round_val[41] = { 0,0,...,0,0,536870912,...,536870912};  
  
int64_t* mult_val_offset = mult_val + 20;  
uint64_t* shift_val_offset = shift_val_array + 20;  
int64_t* round_val_offset = round_val + 20;  
  
int64_t solution7(int64_t mantissa, int16_t from_exp, int16_t to_exp)  
{  
    int16_t diff = to_exp - from_exp;  
    return (mantissa * mult_val_offset[diff] + round_val_offset[diff]) >>  
        shift_val_offset[diff];  
}
```


Solution 7 (15 ns):

```
sub        %esi,%edx  
movswq    %dx,%rdx  
imul      0x41ab60(,%rdx,8),%rdi  
mov       0x41aa00(,%rdx,8),%rcx  
mov       %rdi,%rax  
add       0x41a8a0(,%rdx,8),%rax  
sar       %cl,%rax  
retq
```

- HFT ecosystem has drastically changed over the last decades
- Trading system are required to work with low latency and low jitter
- There are several factors which might impact our system latency and jitter aside from the code
- Measurement is a must, sometimes thing are counter intuitive
- Looking at the assembly which is generated by the compiler helps
- Call for a challenge (www.final.co.il, contactus@final.co.il)





Thank
you!!